

Goal: evaluate self-hosted LLM architectures & agent-loop configurations w.r.t. energy efficiency and task accuracy
(context: autoresearch)

RQ1: configuration effects on energy/accuracy profile

$E_{\text{prop}}, E_{\text{train}}$ (J)

$E_{\text{kept}}, E_{\text{wasted}}$ (J)

val. acc, test acc

RQ2: Pareto-optimal configurations

E_{total} (J)

test acc (final)

wall-clock, kept/reverted counts